

Deduction as a Service

Mohamed Bassem Hassona

German University in Cairo

`mohamed.hassona@student.guc.edu.eg`

Stephan Schulz

DHBW Stuttgart

`schulz@eprover.org`

Abstract

We describe a system offering deduction (over a fixed but flexible background theory) as a service, provided to a client via a network connection. The deduction server maintains the (potentially) large background knowledge base in a form that makes processing of individual conjectures and queries easy and avoids most of the overhead of reading and analyzing large background theories for each individual job. The client connects to the server, can update the background theory by adding and removing client-side or server-side axiom sets, and ask the server to solve proof problems on its behalf.

This offers several benefits: Preprocessing costs can be amortized over multiple proof attempts. The user can be isolated from the complexities of a locally installed ATP system (and indeed locally maintained knowledge bases). Finally, the deduction server can easily offer true strategy parallelism even with sequential back-end theorem provers.

This paper describes the architecture, the communication protocol, and the implementation of a deduction server based on E. First experimental results demonstrate the feasibility of the technology and the performance benefits possible with this approach.

1 Introduction

Classical automated theorem proving has long been concerned with solving individual problems, one at a time. Moreover, problems have often been hand-optimized, either to solve a particular mathematical question, or to analyze and demonstrate the performance of different theorem proving strategies (see e.g. [9]). One of the best-known examples is McCune’s formalization and ultimate proof of the Robbins problem [8].

However, in recent years, theorem provers have been used in a very different setting. Users have developed large theories, either by large-scale manual coding of knowledge as in Cyc [14] or SUMO [12], by organized

Copyright © by the paper’s authors. Copying permitted for private and academic purposes.

In: Pascal Fontaine, Josef Urban, Stephan Schulz (eds.): Proceedings of the 5th Workshop on Practical Aspects of Automated Reasoning (PAAR-2016), Coimbra, Portugal, 02-Jul-2016, published at <http://ceur-ws.org>

efforts to formalize significant parts of mathematics (e.g. MIZAR [20], or more recently Flyspeck [3]), or from the verification of large systems like the sel4 micro-kernel [6]. These theories are typically developed in either an interactive theorem prover, or in authoring environment like Sigma-KEE [13]. Automated theorem provers can be used to dispose a significant number of sub-problems that can be translated into first-order logic. Examples are MizAR for MIZAR and the various “Hammers” for Isabelle [11, 10, 1] and HOL Light/HOL 4 [5, 2]. Communication between interactive and local automatic systems is usually via an ad-hoc protocol. As a fallback, some interactive prover also use the TPTP World web service [18] to invoke remotely installed ATP systems.

In these newer use cases, the theorem prover proves and reproves many different theorems over a large and mostly static background theory. Typically, the background theory has thousands to millions of axioms, only a small fraction of which is used for any given proof. Premise selection techniques [4, 7, 2] enable provers to handle such large problems. However, for large problems, parsing and preprocessing takes significant amounts of time.¹

In this paper, we suggest a different paradigm. Deduction, over a fixed but flexible base theory, is offered as a network service. The deduction server maintains a knowledge base of several pre-loaded background axiomatization. The user can connect to the server, select specific parts of the knowledge base, add additional assumptions, and then request a proof from the server. The server combines the pre-loaded and pre-indexed background theory with the new formulas provided by the user, runs several premise selection strategies to generate ATP problems, and runs several ATP instances with different heuristics to try to solve the problem. If one of the server processes finds a proof, it is handed back to the client.

This approach has a number of advantages. First, the cost of loading and pre-processing the large background theory is amortized over many different proof problems. Proofs to individual problems are often found fast and with low latency. The user can prepare and issue queries from a local desktop, while the deduction server can be shared by several users and run on powerful server hardware. Indeed, the user does not even need to know how to install or invoke ATP systems, since the server can be centrally installed and maintained.

We have developed a suitable communication protocol and implemented a deduction server based on the theorem prover E[15, 16] and its libraries of data types and algorithms for deduction. First results show that the approach is feasible, and that the overhead for processing large problems can be significantly reduced. In the following sections we describe the design, architecture, and implementation of the deduction server, and provide data about a first experimental evaluation.

The system is available at <http://www.e prover.eu/E-eu/DedServer.html>, and will become part of future distributions of E.

2 Client-Server Architecture

2.1 Deduction Server

The deduction server is the central component for offering *Deduction as a Service*. It maintains the current state of the knowledge base, accepts connections from clients, reads and processes their requests, runs deduction jobs on behalf of the clients, and transmits the results back to the client. The main architectural components of the server are the knowledge base, the TCP interface, the axiom filter, and the execution module.

At start-up, the server starts listening on a user-specified port for incoming TCP connections. Whenever a client tries to connect to the server on that port, a process is forked from the main process to serve the client. The server can thus handle several different connections at the same time, and each client is completely isolated from other clients connected to the server.

The client interacts with the server using the protocol specified in section 2.3. The server implements a typical read-execute-print loop, accepting commands from the client, executing them, and sending the results back to the client. The client can upload axiom sets, remove them, and query the server for proofs and answers (i.e. instantiations of variables that make an existentially quantified query formula true [19]).

Whenever the client uploads a set of axioms, the server parses the axioms, adds them to various indexes useful for axiom selection/pruning methods, and adds them to the in-memory knowledge base. However, these uploaded axiom sets will not automatically be used by each proof attempt, but only when the client has also activated (or *staged*) the particular set in the current session. Thus, the server maintains a pre-indexed, ready-to-use library of axiom sets. In addition to axiom sets uploaded by the client, the server can also offer access to axiom sets

¹As an example, E takes around 54 seconds (on a system with a 4 GHz Intel i7 and a fast SSD) just to parse the TPTP problem `CSR073+6.p`, a first-order problem based on OpenCyc with nearly 3.5 million axioms and taking up 480 MB. Indexing for SInE axiom selection takes a further 2 seconds, while more than half of successful proof attempts need less than 1 second.

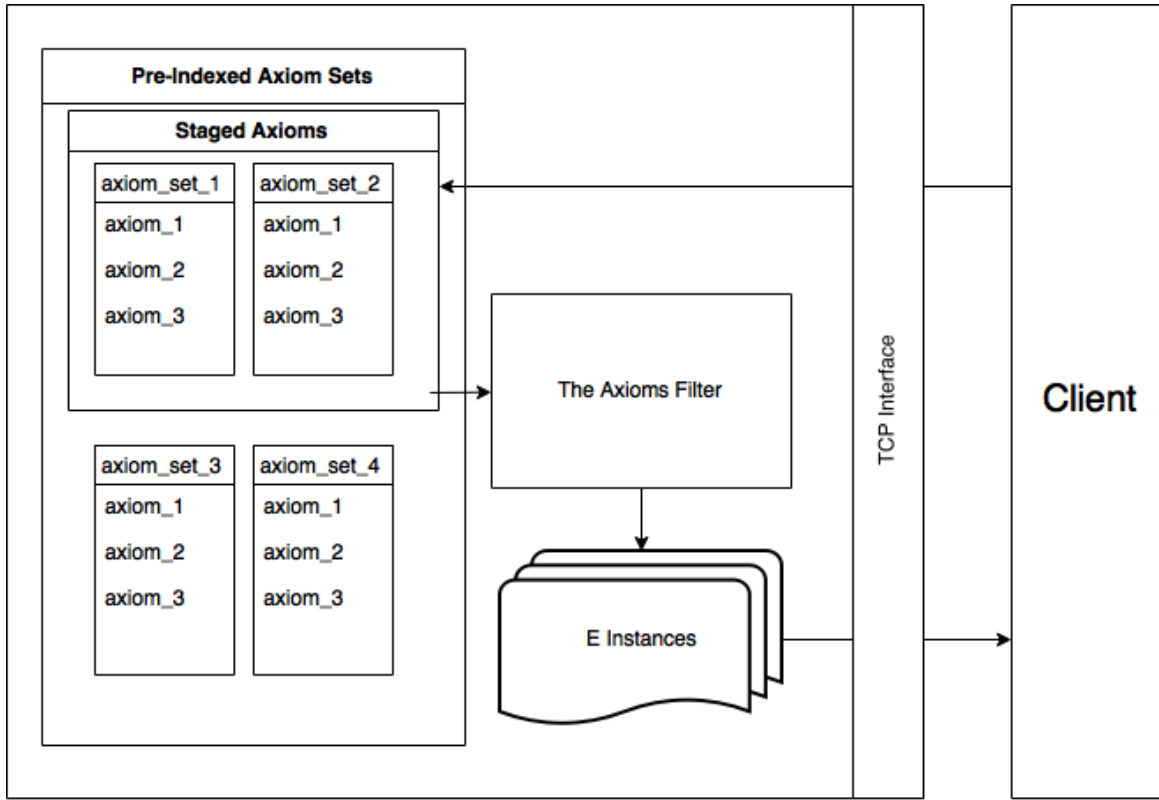


Figure 1: The Deduction Server Architecture

stored in a server-side library on disk. Axioms loaded by the server on start-up are available to all users, while axiom sets added by a user during a session are currently only available in that session.

After the client chooses some axiom sets to be used in the proof search, it can start querying the server. The query consists of optional query-specific axioms and a conjecture. These additional formulas are discarded immediately after the query has been processed. The server adds the new axiom set to the currently staged part of the knowledge base and hands the problem to the axiom filter. The axiom filter module applies one or multiple different relevancy pruning strategies to the extended knowledge base. Each of the pruning strategies produces one much smaller proof problem. The pruned problems are passed to the execution module, which starts several instances of E in parallel and monitors their success or failure. Whenever one of the running E instances finds a proof for the conjecture, the other instances are stopped and the proof is returned back to the user. Currently, all instances of E use its conventional *automatic mode*, and rely on the different pruning strategies to provide diversity for the proof attempt.

The rough architecture of the deduction server is shown in Figure 1.

The deduction server is written in C and integrated with the E distribution. It uses the E libraries and data types to parse and represent logical formulas, to maintain and process first order knowledge bases, and to perform axiom selection.

2.2 User client

The E distribution includes a small Python client called “enetcat.py”, which can be used to interact with the server. The client is intended as a reference implementation, to demonstrate the interactions with the server and enable other developers to use it as a model for interfacing the server with other systems.

The client requires the address and the port of the deduction server as arguments. It opens a TCP connection to the server and then reads the user’s commands from “stdin”, communicates them to the server, and prints the output back to “stdout”.

In addition to encoding the plain text commands into the distinct length-delimited TCP messages expected by the server, the client offers two convenience features. First, it offers command line editing. Secondly, it locally expands TPTP style `include` commands for formula sets to be uploaded to the server. This makes uploading

of even large axioms sets plausible and convenient.

The current client is tailored towards command-line users. However, the protocol is simple enough to be implemented in a few lines of code in nearly any modern programming language, making other clients, in particular a web-based client, easily possible.

2.3 Communication Protocol

As mentioned before, the communication between the server and the client is done over TCP to ensure reliability. TCP messages between the server and the client are encoded in a way to ease the communication between them. Each TCP messages starts with 4 bytes containing the length of the messages, including those 4 bytes, followed by the actual message. The commands themselves are plain, human-readable ASCII text.

We describe the protocol for communication between the server and the client using a running example. We assume that the connection has already been established and all messages are encoded as described in the previous paragraph.

The first section of the session uploads two small axiom sets to the server.

```
1 ADD axiom_set1
2 fof(inp1,axiom,(subclass(a,b))).
3 fof(inp2,axiom,(subclass(b,c))).
4 GO
5 >205 ok : added
6
7 ADD axiom_set2
8 fof(inp3,axiom,((subclass(X,Y) & subclass(Y,Z)) => subclass(X,Z))).
9 GO
10 >205 ok : added
```

Each of the “ADD ...GO” blocks upload one named set of axioms to the system. Axioms uploaded are parsed and stored in the memory of the server, but are not automatically included in future proof attempts. The server responds with the success status of each command’s execution.

```
1 STAGE axiom_set1
2 >201 ok : staged
3
4 LIST
5 >Staged :
6 > axiom_set1
7 >Unstaged :
8 > axiom_set2
9 >On Disk :
10 > test.p
11 >200 ok : success
12
13 STAGE axiom_set2
14 >201 ok : staged
```

To select axiom sets for future proof attempts, we use the “STAGE” command. It marks the named axiom set as active for later proof attempts. The “LIST” command provides the status of all axiom sets currently known to the server. Possible states are *Staged*, *Unstaged* (but in memory and pre-indexed for axiom selection), and *On Disk*, i.e. known to the server, but not loaded or indexed. Server-side axiom sets can be loaded into the server’s memory using the “LOAD” command.

```
1 RUN job1
2 fof(inp4,conjecture,(subclass(a,c))).
3 GO
4 # Pid: 27988
5 # Auto-Mode selected heuristic G_E__208_C18_F1_SE_CS_SP_PS_S0Y
6 # and selection function SelectMaxLComplexAvoidPosPred.
7 #
```

```

8 # Preprocessing time      : 0.010 s
9 # Presaturation interreduction done
10
11 # Proof found!
12 # SZS status Theorem
13 # SZS output start CNFRefutation.
14
15 fof(c_0_0, axiom,
16 (((subclass(X1,X2)&subclass(X2,X3))=>subclass(X1,X3))),
17 file('/var/folders/__/ss_kh09s5_l9s1twdz7k5y900000gn/T//epr_VKSv9K',
18 i_0_3)).
19 fof(c_0_1, conjecture, (subclass(a,c)),
20 file('/var/folders/__/ss_kh09s5_l9s1twdz7k5y900000gn/T//epr_VKSv9K',
21 i_0_4)).
22 fof(c_0_2, axiom, (subclass(b,c)),
23 file('/var/folders/__/ss_kh09s5_l9s1twdz7k5y900000gn/T//epr_VKSv9K',
24 i_0_2)).
25 fof(c_0_3, axiom, (subclass(a,b)),
26 file('/var/folders/__/ss_kh09s5_l9s1twdz7k5y900000gn/T//epr_VKSv9K',
27 i_0_1)).
28 fof(c_0_4, plain,
29 (((~subclass(X1,X2)|~subclass(X2,X3))|subclass(X1,X3))),
30 inference(fof_nnf,[status(thm)],[c_0_0])).
31 fof(c_0_5, negated_conjecture, (~subclass(a,c)),
32 inference(fof_simplification,[status(thm)],[inference(assume_negation,
33 [status(cth)],[c_0_1])))).
34 cnf(c_0_6, plain, (subclass(X1,X2)|~subclass(X3,X2)|~subclass(X1,X3)),
35 inference(split_conjunct,[status(thm)],[c_0_4])).
36 cnf(c_0_7, plain, (subclass(b,c)), inference(split_conjunct,
37 [status(thm)],[c_0_2])).
38 cnf(c_0_8, negated_conjecture, (~subclass(a,c)),
39 inference(split_conjunct,[status(thm)],[c_0_5])).
40 cnf(c_0_9, plain, (subclass(X1,c)|~subclass(X1,b)),
41 inference(spm,[status(thm)],[c_0_6, c_0_7])).
42 cnf(c_0_10, plain, (subclass(a,b)), inference(split_conjunct,
43 [status(thm)],[c_0_3])).
44 cnf(c_0_11, negated_conjecture, ($false), inference(cn,[status(thm)],
45 [inference(rw,[status(thm)],[inference(spm,[status(thm)],[c_0_8,
46 c_0_9]),
47 c_0_10]))], ['proof'])).
48 # SZS output end CNFRefutation.
49
50 # -----
51 # User time          : 0.008 s
52 # System time        : 0.002 s
53 # Total time         : 0.010 s
54 # Maximum resident set size: 2838528 pages
55
56 # Processing finished for job1
57
58 200 ok : success

```

After staging the needed axioms, we can start running a job with “RUN ...GO” block. Formulas introduced in the RUN command block are used only temporally for this particular job, and are discarded after its termination. These formulas typically contain the conjecture, but can also include additional axioms and assumptions. If the server succeeds in proving the conjecture, it provides back the logical status of the query, and can also include a proof object.

```
1 UNSTAGE axiom_set2
```

```

CSR025+6.p CSR026+6.p CSR027+6.p CSR028+6.p CSR029+6.p CSR030+6.p
CSR031+6.p CSR032+6.p CSR033+6.p CSR034+6.p CSR035+6.p CSR036+6.p
CSR037+6.p CSR038+6.p CSR039+6.p CSR040+6.p CSR041+6.p CSR042+6.p
CSR043+6.p CSR044+6.p CSR045+6.p CSR046+6.p CSR047+6.p CSR048+6.p
CSR049+6.p CSR050+6.p CSR051+6.p CSR052+6.p CSR053+6.p CSR054+6.p
CSR055+6.p CSR056+6.p CSR057+6.p CSR058+6.p CSR059+6.p CSR060+6.p
CSR061+6.p CSR062+6.p CSR063+6.p CSR064+6.p CSR065+6.p CSR066+6.p
CSR067+6.p CSR068+6.p CSR069+6.p CSR070+6.p CSR071+6.p CSR072+6.p
CSR073+6.p CSR074+6.p CSR111+6.p

```

Figure 2: List of problems used in the evaluation

```

2  >202 ok : unstaged
3
4  REMOVE axiom_set2
5  >203 ok : removed
6
7  DOWNLOAD axiom_set1
8  >fof(inp1,axiom,(subclass(a,b))).
9  >fof(inp2,axiom,(subclass(b,c))).
10 >204 ok : downloaded
11
12 RUN job2
13 fof(inp4,conjecture,(subclass(a,c))).
14 GO
15 # Processing started for job2
16 # SZS status GaveUp for job2
17
18 # Processing finished for job2
19
20 200 ok : success

```

Axiom sets can also be “UNSTAGE”d or “REMOVE”d completely from the system. Axiom sets can be retrieved from the server using the “DOWNLOAD” command. The example shows how we remove one of the axiom sets and re-run the job. Without the necessary axioms, the proof is not found.

3 Evaluation

For the experimental evaluation, we selected a set of large problems with a common axiomatization from the TPTP library [17], version 6.0.0. In particular, we used all problems that include the axiom set `CSR002+5.ax`, a translation of the OpenCyc knowledge base into first-order logic [14]. The axiom set contains 3 341 983 axioms and occupies about 479 megabytes disk space in TPTP format. It is used by 51 problems, each of which combines it with a single unique conjecture. See Figure 2 for the list of problems.

To illustrate the comparative performance of the service model and the conventional one-problem-at-a-time model, we configured both versions to run exactly the same problems with exactly the same pruning parameters and E’s sequential automatic mode. In particular, the server did not employ any strategy parallelism in pruning or search, but rather selected the same search strategy than the stand-alone prover for each a given problem.

The systems were given a time limit of 30 seconds for actual proof search, in addition to the time taken for parsing and axiom selection for the standalone version (about 110 seconds per problem on this hardware, with some variation). Memory was limited to 1024MB. Tests were performed on a virtualized server with 8 2.6 GHz Intel CPUs, of which only one was effectively utilized by our test setup. With this configuration, both setups solved 41 of the 51 problems.

We recorded the time taken after 5, 10, 20, 30, 40 and all 51 problems.

The results in Figure 3 show the difference between the single strategy server mode and the conventional prover. The figure shows the accumulated wall-clock time over the number of problems for which a proof was attempted. In the conventional case, the prover has to parse and index the axiom set for each proof attempt. In the deduction server, the axioms are parsed and indexed for axiom selection only once. The are then kept in




Figure 3: Benchmarking Results

the server’s memory and multiple queries can be executed against this axiom set. Thus, the pre-processing time of each problem in the single strategy server mode is amortized over multiple runs.

4 Conclusion

We have described the first implementation of a Deduction Server and have introduced a communication protocol that supports remote reasoning over largely fixed but flexible domains as a service. Our results shows that the deduction server, even if configured for identical pruning and search strategies, improves the run time for sets of related problems by a significant amount, due to the amortization of costs for parsing and the indexing needed for axiom selection. This clearly shows that the approach is feasible and has promise even if looking at it only from a performance point of view.

In addition to these performance benefits, having a formal protocol to communicate with the prover over the net makes integrating it with other applications much easier. Having access to a remotely running prover is a win for those who can’t install provers locally, e.g. for OS compatibility issues or due to insufficient local computing resources.

Future work includes improved multi-user support, in particular controlling both access, but also the resources, such as CPU time and memory, allocated for each user. Another important step would be the integration of multiple deduction systems in the back-end, to improve overall performance. This would offer a single interface to potentially very diverse deduction systems.

Finally, the Deduction Server can be extended to a cluster of servers, to offer real scalable deduction as a service. In the simplest version, a single head node accepts and processes user commands, pre-processes the problems, and distributes the deduction jobs to different servers. An even more scalable version would maintain different user sessions with the current axiomatization on dedicated machines, so that axiom selection can also be distributed.

References

- [1] Sascha Böhme and Tobias Nipkow. Sledgehammer: Judgement Day. In Jürgen Giesel and Reiner Hähnle, editors, *Proc. of the 5th IJCAR, Edinburgh*, volume 6173 of *LNAI*, pages 107–121. Springer, 2010.

- [2] Thibault Gauthier and Cezary Kaliszyk. Premise Selection and External Provers for HOL4. In *Proceedings of the 2015 Conference on Certified Programs and Proofs*, pages 49–57, New York, NY, USA, 2015. ACM.
- [3] Thomas C. Hales, John Harrison, Sean McLaughlin, Tobias Nipkow, Steven Obua, and Roland Zumkeller. A revision of the proof of the Kepler conjecture. In Jeffrey C. Lagarias, editor, *The Kepler Conjecture*, pages 341–376. Springer, 2011.
- [4] Kryštof Hoder and Andrei Voronkov. Sine Qua Non for Large Theory Reasoning. In Nikolaj Bjørner and Viorica Sofronie-Stokkermans, editors, *Proc. of the 23rd CADE, Wroclaw*, volume 6803 of *LNAI*, pages 299–314. Springer, 2011.
- [5] Cezary Kaliszyk and Josef Urban. Learning-assisted automated reasoning with Flyspeck. *Journal of Automated Reasoning*, 53(2):173–213, 2014.
- [6] Gerwin Klein, Kevin Elphinstone, Gernot Heiser, June Andronick, David Cock, Philip Derrin, Dhammika Elkaduwe, Kai Engelhardt, Rafal Kolanski, Michael Norrish, Thomas Sewell, Harvey Tuch, and Simon Winwood. seL4: Formal verification of an OS kernel. In *Proc. of the 22nd ACM Symposium on Principles of Operating Systems (SOSP)*, Big Sky, Montana, USA, October 2009. ACM.
- [7] Daniel Kühlwein, Twan van Laarhoven, Evgeni Tsivtsivadze, Josef Urban, and Tom Heskes. Overview and evaluation of premise selection techniques for large theory mathematics. In Bernhard Gramlich, Ulrike Sattler, and Dale Miller, editors, *Proc. of the 6th IJCAR, Manchester*, volume 7364 of *LNAI*, pages 378–392. Springer, 2012.
- [8] William McCune. Solution of the Robbins problem. *Journal of Automated Reasoning*, 19(3):263–276, 1997.
- [9] W.W. McCune. 33 Basic Test Problems: A Practical Evaluation of Some Paramodulation Strategies. In R. Veroff, editor, *Automated Reasoning and its Applications: Essays in Honor of Larry Wos*, chapter 5, pages 71–114. MIT Press, 1997.
- [10] Jia Meng and Lawrence C. Paulson. Translating higher-order clauses to first-order clauses. *Journal of Automated Reasoning*, 40(1):35–60, 2008.
- [11] Jia Meng and Lawrence C. Paulson. Lightweight relevance filtering for machine-generated resolution problems. *Journal of Applied Logics*, 7(1):41–57, 2009.
- [12] Adam Pease, Ian Niles, and John Li. The Suggested Upper Merged Ontology: A Large Ontology for the Semantic Web and its Applications. In *Working Notes of the AAAI-2002 Workshop on Ontologies and the Semantic Web*, 2002.
- [13] Adam Pease and Stephan Schulz. Knowledge Engineering for Large Ontologies with Sigma KEE 3.0. In Stephane Demri, Deepak Kapur, and Christoph Weidenbach, editors, *Proc. of the 7th IJCAR, Vienna*, volume 8562 of *LNAI*, pages 519–525. Springer, 2014.
- [14] Deepak Ramachandran, Pace Reagan, and Keith Goolsbey. First-orderized ResearchCyc: Expressiveness and Efficiency in a Common Sense Knowledge Base. In Pavel Shvaiko, editor, *Proc. of the AAAI Workshop on Contexts and Ontologies: Theory, Practice and Applications (C&O-2005)*, 2005.
- [15] S. Schulz. E – A Brainiac Theorem Prover. *Journal of AI Communications*, 15(2/3):111–126, 2002.
- [16] Stephan Schulz. System Description: E 1.8. In Ken McMillan, Aart Middeldorp, and Andrei Voronkov, editors, *Proc. of the 19th LPAR, Stellenbosch*, volume 8312 of *LNCS*. Springer, 2013.
- [17] Geoff Sutcliffe. The TPTP Problem Library and Associated Infrastructure: The FOF and CNF Parts, v3.5.0. *Journal of Automated Reasoning*, 43(4):337–362, 2009.
- [18] Geoff Sutcliffe. The TPTP World - infrastructure for automated reasoning. In E. Clarke and A. Voronkov, editors, *Proc. of the 16th LPAR, Dakar*, number 6355 in *LNAI*, pages 1–12. Springer, 2010.
- [19] Geoff Sutcliffe, Mark Stickel, Stephan Schulz, and Josef Urban. Answer Extraction for TPTP. <http://www.cs.miami.edu/~tptp/TPTP/Proposals/AnswerExtraction.html>. (accessed 2013-07-08).

- [20] Andrzej Trybulec and Howard A. Blair. Computer assisted reasoning with MIZAR. In *Proc. of the 9th IJCAI, Los Angeles*, volume 85, pages 26–28. Morgan Kaufmann, 1985.